

# Password hashing

- Introduced in 1970s
- Sense of security which is often false
- Fast hashes
  - ▶ MD4, MD5, SHA-1
  - ▶ Designed for MACs and digital signatures
  - ▶ Must be easy to compute
- Slow hashes
  - ▶ bcrypt, PBKDF2, scrypt
  - ▶ Designed for password hashing
  - ▶ Inefficient and difficult to compute

način rada	cijena	sol	sažetak
\$2y\$	10	\$VaohDzrEa32G3DuTqduoCe	05cwi9EsLCniXVABkXeC2xq4z74ZbYG


# Motivation

- Bcrypt is:
  - ▶ Popular (often used): WWW, OpenBSD, SUSE Linux
  - ▶ Slow
  - ▶ Sequential
  - ▶ Random memory footprint
  - ▶ Designed to be resistant to brute force attacks and to remain secure despite hardware improvements
- You could almost think why even bother optimizing

# Bcrypt (Provos and Mezieres)

- Based on Blowfish block cipher (Bruce Schneier)
- Expensive key setup
- User defined cost setting
  - ▶ Cost setting between 4 and 31 inclusive is supported
  - ▶ Cost 5 is traditionally used for benchmarks for historical reasons
  - ▶ All given performance figures are for bcrypt at cost 5
  - ▶ Current systems should use higher cost setting
- Pseudorandom memory accesses
- Memory usage
  - ▶ 4 KB for four S-boxes
  - ▶ 72 B for P-box

# Blowfish encryption



- Symmetric-key cipher
- 64-bit input block, 32 to 448-bit key
- Feistel network
- Pseudorandom memory accesses
  - ▶ 32-bit loads from four 1 KB S-boxes initialized with digits of number  $\pi$

$$R_i = L_{i-1} \oplus P_i \quad (1)$$

$$L_i = R_{i-1} \oplus F(R_i) \quad (2)$$

$$F(a, b, c, d) = ((S_1[a] + S_2[b]) \oplus S_3[c]) + S_4[d] \quad (3)$$

Niels Provos and David Mazieres, "A Future-Adaptable Password Scheme", The OpenBSD Project, 1999

# EksBlowfish

## Ekspensive key schedule Blowfish

---

**Algorithm 1** EksBlowfishSetup(cost, salt, key)

---

```
1:  $state \leftarrow \text{InitState}()$ 
2:  $state \leftarrow \text{ExpandKey}(state, salt, key)$ 
3:  $\text{repeat}(2^{\text{cost}})$ 
4:    $state \leftarrow \text{ExpandKey}(state, 0, salt)$ 
5:    $state \leftarrow \text{ExpandKey}(state, 0, key)$ 
6:  $\text{return } state$ 
```

---

- Order of lines 4 and 5 is swapped in implementation

# bcrypt

---

**Algorithm 2** `bcrypt(cost, salt, pwd)`

---

- 1:  $state \leftarrow \text{EksBlowfishSetup}(cost, salt, key)$
  - 2:  $ctext \leftarrow \text{"OrpheanBeholderScryDoubt"}$
  - 3:  $\text{repeat}(64)$
  - 4:      $ctext \leftarrow \text{EncryptECB}(state, ctext)$
  - 5:  $\text{return Concatenate}(cost, salt, ctext)$
- 

Output: (bcrypt\_indicator, cost, 128-bit base-64 22 chars salt, 184-bit base-64 31 chars hash)

\$2a\$12\$GhvMmNVjRW29ulnudl.LbuAnUtN/LRfe1JsBm1Xu6LE3059z5Tr8m

# Architecture

## Epiphany

- 16/64 32-bit RISC cores operating at up to 1 GHz/800 MHz
  - ▶ Chips used in testing operate at 600 MHz
- Pros
  - ▶ **Energy-efficient** - 2 W maximum chip power consumption
  - ▶ 32 KB of local memory per core
  - ▶ 64 registers
  - ▶ FPU can be switched to integer mode
  - ▶ Dual-register (64-bit) load/store instructions
- Cons
  - ▶ FPU in integer mode can issue only add and mul instructions
  - ▶ Only simple addressing modes
    - Index scaling would be helpful for S-box lookups

# Implementation

## Epiphany

- John the Ripper prepares data on ARM cores
- Bcrypt hashes computed on Epiphany
- Single instance does not have enough instruction level parallelism
  - ▶ FPU has four cycle latency
  - ▶ FPU does not have bitwise instructions
- Optimized in assembly
- Two instances overlapped to exploit dual-issue architecture
  - ▶ Integer ALU
  - ▶ FPU in integer mode



# Architecture

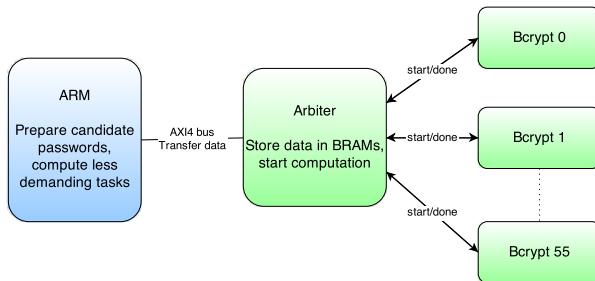
## Zynq 7020

- Heterogeneous device
- Dual ARM Cortex-A9 MPCore
  - ▶ 667 MHz
  - ▶ 256 KB on-chip memory
- Advanced low power 28nm programmable logic
  - ▶ 85 K logic cells
  - ▶ 560 KB of block RAM
- AXI buses used for CPU-FPGA communication

# Implementation

## Zynq 7020

- John the Ripper prepares data on ARM cores
- Bcrypt instances compute hash
- Number of concurrent instances limited by available BRAM
- Large communication overhead for low cost setting
- Hardware defects of ZedBoard limit optimizations



# Zynq 7045

- Architecture
  - ▶ ARM CPU and Zynq reconfigurable logic
  - ▶ Roughly 4 times bigger than Zynq 7020
- Implementation
  - ▶ ZedBoard implementation ported to a bigger device
  - ▶ Bottleneck: CPU-FPGA communication (for low cost setting)
  - ▶ Not possible to use all available resources due to hardware defects of ZC706 board

# Theoretical Peak Performance Analysis

## Theory

$$c/s = \frac{N_{ports} * f}{(2^{cost} * 1024 + 585) * N_{reads} * 16} \quad (4)$$

- $N_{ports}$  - number of available read ports to local memory or L1 cache
- $N_{reads}$  - number of reads per Blowfish round
  - ▶ 4 or 5 depending on whether reads from P-boxes go from one of those read ports we've counted or from separate storage such as registers
- $2^{cost} * 1024 + 585$  - number of Blowfish block encryptions in bcrypt hash computation
- $f$  (in Hz) - clock rate

---

```
bcrypt(cost, salt, pwd)
```

---

```

1: state ← InitState()
2: state ← ExpandKey(state, salt, key)
3: repeat( $2^{cost}$ )
4:   state ← ExpandKey(state, 0, salt)
5:   state ← ExpandKey(state, 0, key)
6: ctext ← "OrpheanBeholderScryDoubt"
7: repeat(64)
8:   ctext ← EncryptECB(state, ctext)
9: return Concatenate(cost, salt, ctext)
```

---

# Takeaways

- Many-core low power RISC platforms and FPGAs are capable of exploiting bcrypt peculiarities to achieve comparable performance and higher energy-efficiency
- Higher energy-efficiency enables higher density
  - ▶ More chips per board, more boards per system
- It doesn't take ASICs to improve bcrypt cracking energy-efficiency by a factor of 45+
  - ▶ Although ASICs would do better yet